

# trafficCrash

Frank Guo

2024-04-15

## Milestone 2 - Crashes Analysis

By Xiaodong Guo

### 1. Goal.

Our project has a significant objective- to construct models identifying the pivotal factors contributing to severe crashes. This crucial task is based on the data provided by the New Zealand Transport Agency(NZTA).

### 2. Data Source.

The original data, meticulously collected, originated from the Waka Kotahi NZ Transport Agency's open data portal(the tutor provided the link in the assignment piece). We specifically downloaded the dataset named "Crash Analysis System (CAS) data" from the "Crash" catalogue, which encompasses all traffic crashes reported to us by the NZ Police. The data format is a "CSV" file. It was created on 3/25/2020 and last updated on 3/14/2024.

The data includes crash datas from 2000 to 2023.

### 3. Data Processing.

```
#load data

library(tidyverse)
library(xgboost)

#file_path <- file.choose()
set.seed(230)
data <- read.csv("../data/Crash_Analysis_System_(CAS)_data.csv", header = TRUE, sep = ",")

dim(data)

## [1] 821744      72

#glimpse(data)
```

We load the data from csv file. The dataset we got have 72 columns, and 821744 rows. The first things we can do is to drop columns not related to our objective. Thus, we select following columns by common sense of mine (not sure if 100% correct, maybe need some hint from Lisa Chen). There are also have some description columns, that is long string values to describe an event or street name. That looks no sense, should drop them too. Like "crashLocation1", "crashLocation2". Then drop columns with almost all values are Null, that is "crashRoadSideRoad" "intersection".

Define crashSeverity == "Fatal Crash" and crashSeverity == "Serious Crash" as severe crashes given numeric value 1,

Define crashSeverity == "Minor Crash" | crashSeverity == "Non-Injury Crash" as not severe crashes given numeric value 0.

The "crashSeverity" Label will be the target label.

Predictors like "weatherA" and "weatherB", could be combined as one attribute "weather", much easier to display and dealing with it later.

Replace the "", "Null", "None" value in region with "Other", replace the "" in other character attributes with "others".

List all the attributes with Na value: "advisorySpeed" "bicycle" "bridge" "bus" "carStationWagon" "cliffBank" "ditch" "fence" "guardRail" "houseOrBuilding" "kerb" "moped" "motorcycle" "NumberOfLanes" "otherVehicleType" "overBank" "parkedVehicle" "pedestrian" "postOrPole" "roadworks" "schoolBus" "slipOrFlood" "speedLimit" "strayAnimal" "suv" "taxi" "temporarySpeedLimit" "trafficIsland" "trafficSign" "train" "tree" "truck" "unknownVehicleType" "vanOrUtility" "vehicle" "waterRiver"

They are all numeric, and fill with 0 is reasonable.

The whole dataset looks categorical. So we convert all attributes in to categorical attributes.

```
#clean data

columns_to_drop <- c("X", "Y", "OBJECTID", "areaUnitID", "crashDirectionDescription", "", "crashDistance", "tl",
"crashFinancialYear", "fatalCount", "debris", "meshblockId", "northing", "easting", "objectThrownOrDropped", "l",
"otherObject", "phoneBoxEtc", "seriousInjuryCount")

data <- select(data, -one_of(columns_to_drop))

na_percentage <- colMeans(is.na(data))
columns_with_high_na <- names(na_percentage[na_percentage > 0.99])
#print(columns_with_high_na)

data <- data %>% select(-columns_with_high_na)

#table(data$crashSeverity)

data <- data %>%
  mutate(crashSeverity = ifelse(crashSeverity == "Fatal Crash" | crashSeverity == "Serious Crash", 1,
    ifelse(crashSeverity == "Minor Crash" | crashSeverity == "Non-Injury Crash", 0,
      2))) %>% filter(crashSeverity != 2)

#table(data$crashSeverity)

#table(data$weatherA)
#table(data$weatherB)

data$weatherA <- ifelse(data$weatherA %in% c("None", "Null"), "Other", data$weatherA)
data$weatherB <- ifelse(data$weatherB %in% c("None", "Null"), "", data$weatherB)

data <- data %>% unite(weatherA, weatherB, col=weather, sep=" ")

table(data$region)
```

```
##
##           Auckland Region      Bay of Plenty Region
##           3188                285346                47177
##           Canterbury Region    Gisborne Region      Hawke's Bay Region
##           82146                9784                 32388
##           Manawatū-Whanganui Region Marlborough Region Nelson Region
##           46329                8266                 8076
##           Northland Region      Otago Region        Southland Region
##           33299                44574                 20234
##           Taranaki Region       Tasman Region      Waikato Region
##           18604                7541                 87849
##           Wellington Region     West Coast Region
##           79725                7218

data$region <- ifelse(data$region %in% c("None", "Null"), "Other", data$region)

na_columns <- sapply(data, function(x) any(is.na(x)))

columns_with_na <- names(data)[na_columns]
#print(columns_with_na)

data <- data %>%
  mutate_at(vars(one_of(columns_with_na)), ~replace_na(., 0))

processed_data <- data
processed_data[processed_data == ""] <- "others"

processed_data[] <- lapply(processed_data, function(x) as.factor(x))
data <- processed_data

rm(processed_data)

#glimpse(data)
```

## 4. Data Exploration

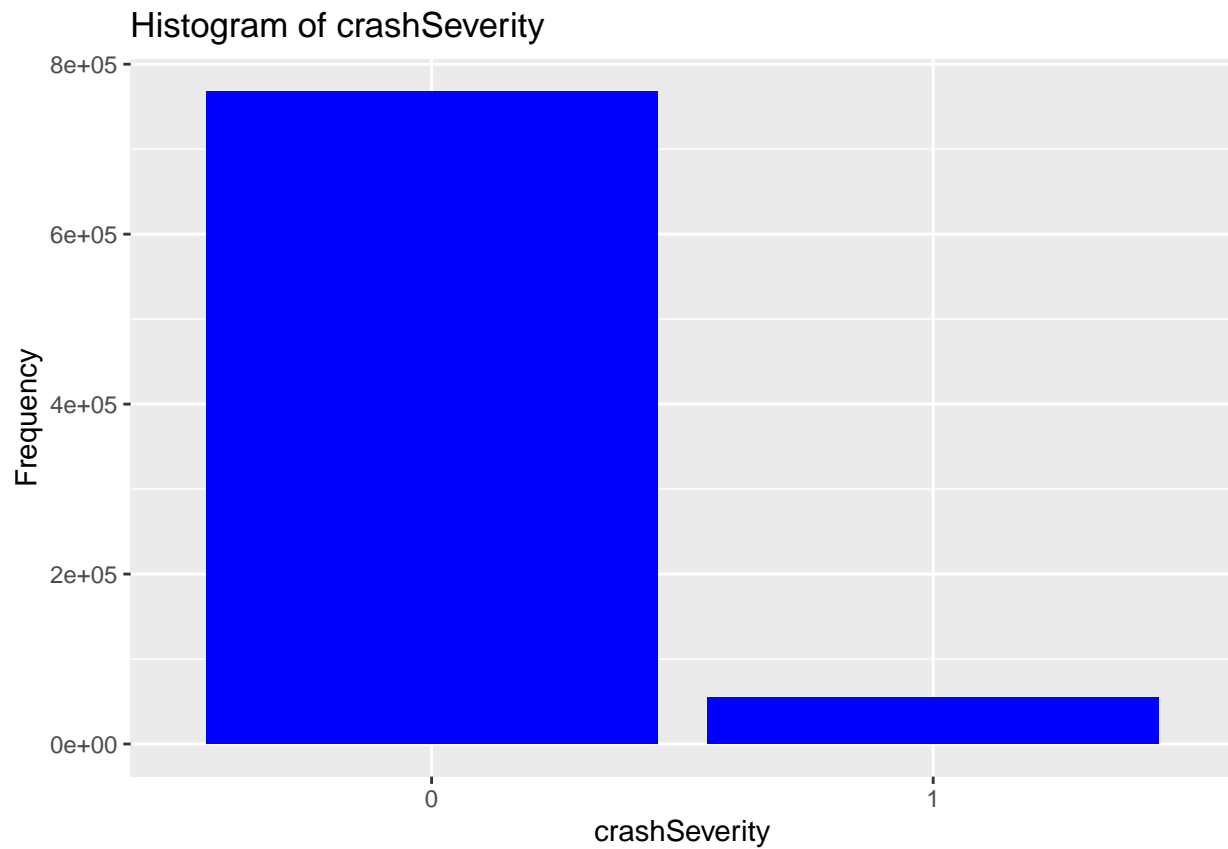
Because all the attributes are categorical attributes, we decide to explore the frequency the most important attributes.

We Plot the target labels and some relations between attributes like crashYear, region, weather and light which is intuitively think that will be key attribute of the cause of incident.

We Can see total crashes decreased by year, but severe crashes not. Also, we found the problem that unbalanced data occurs in predictors. For examples, regarding as weather, we can see most of the crashes (severe and regular) happened in fine day rather than the other weather conditions.

From All of those plots, we found the unbalance of the data. There are 767290 regular crashes, but only 54454 severe crashes. Our target is to find the cause of severe crashes, but the severe observations' size is really small compares to the regular crashes.

Var1	Freq
0	767290
1	54454

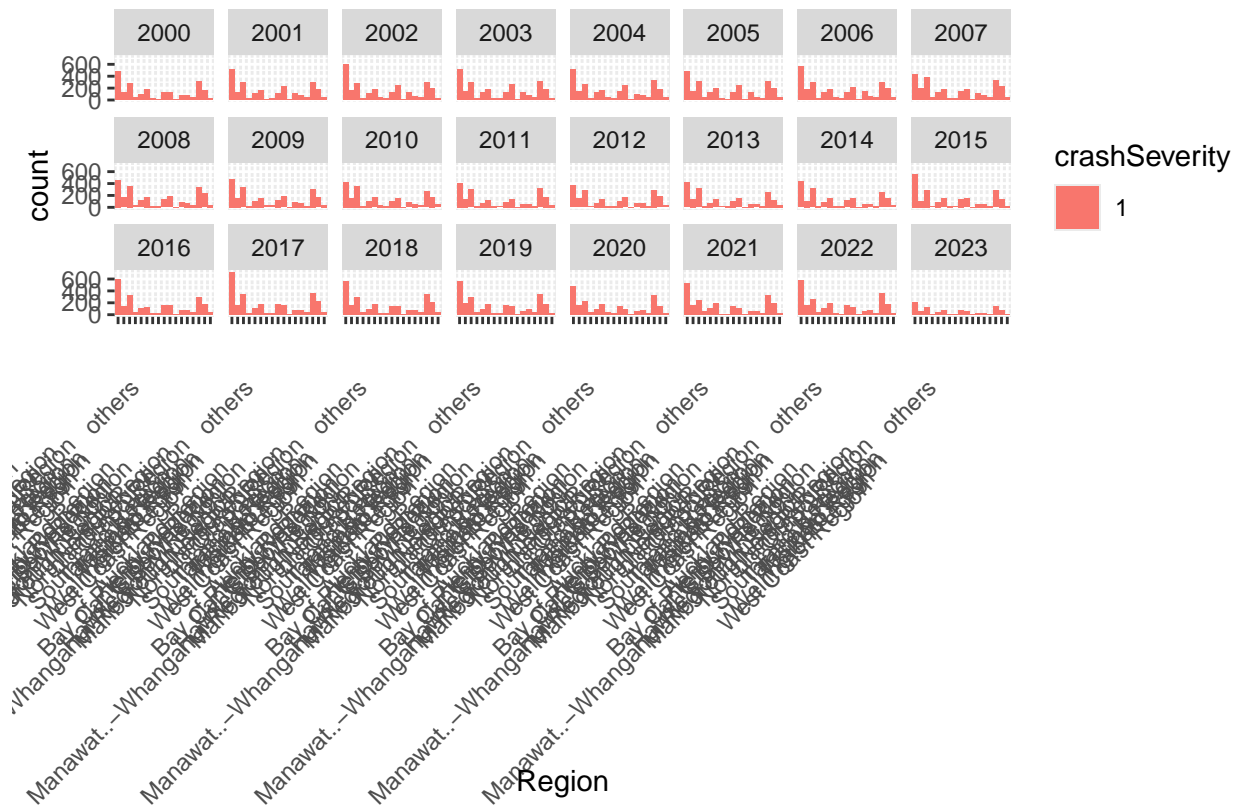


```
## `summarise()` has grouped output by 'crashYear', 'region'. You can override  
## using the `.groups` argument.
```


Comparison of Severe and Regular Values by Region



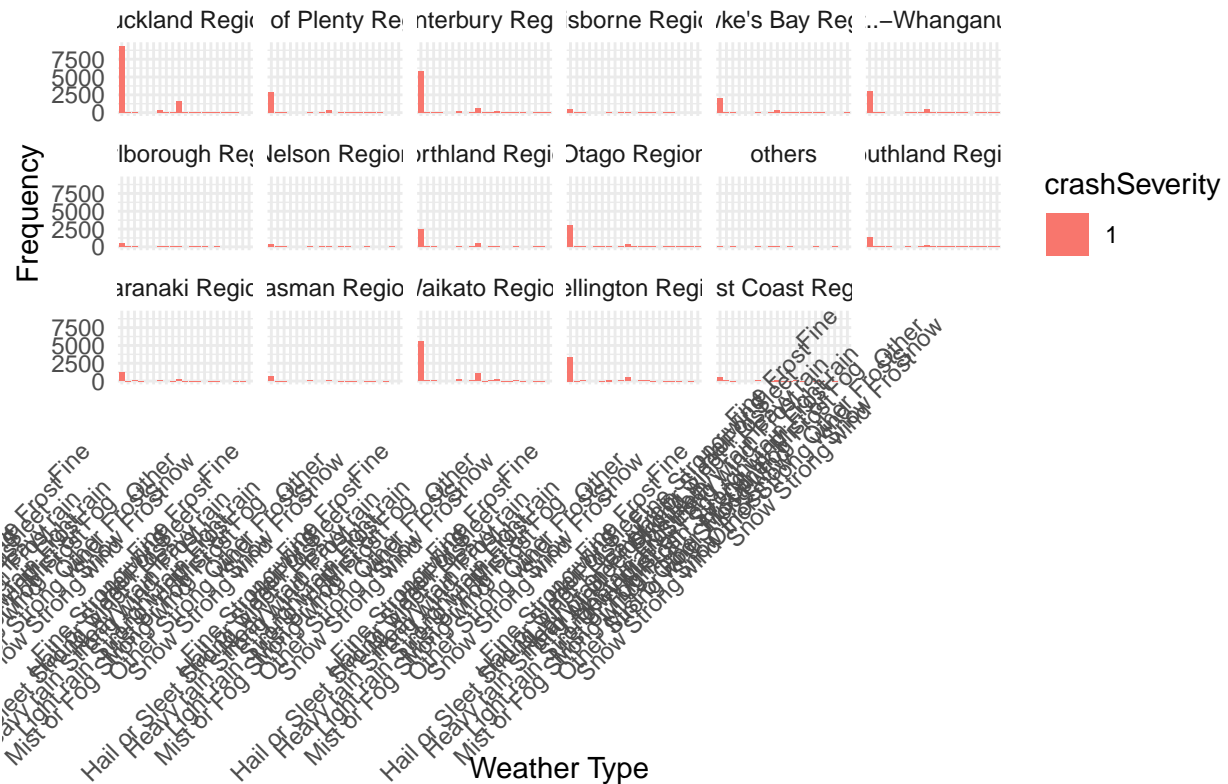
Comparison of Severe and Regular Values by Region



```
## `summarise()` has grouped output by 'region', 'weather'. You can override using
## the `.groups` argument.
```

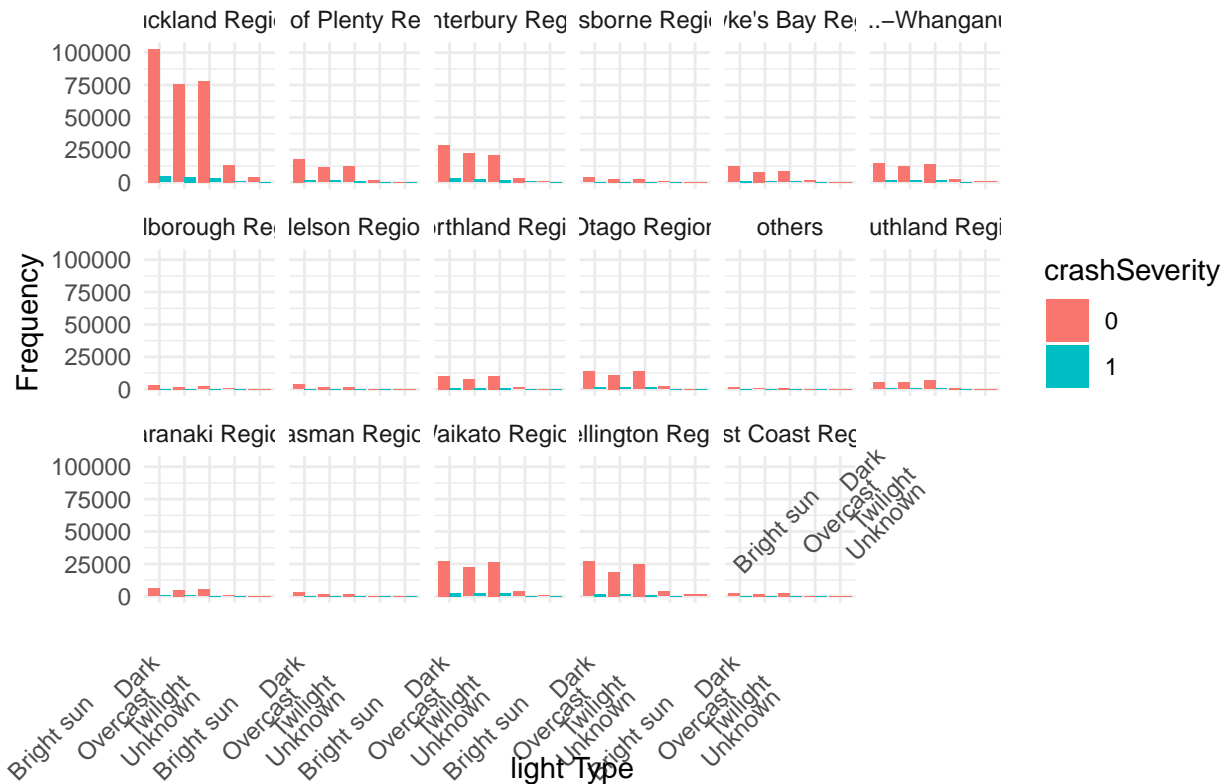


## Weather Conditions in Region

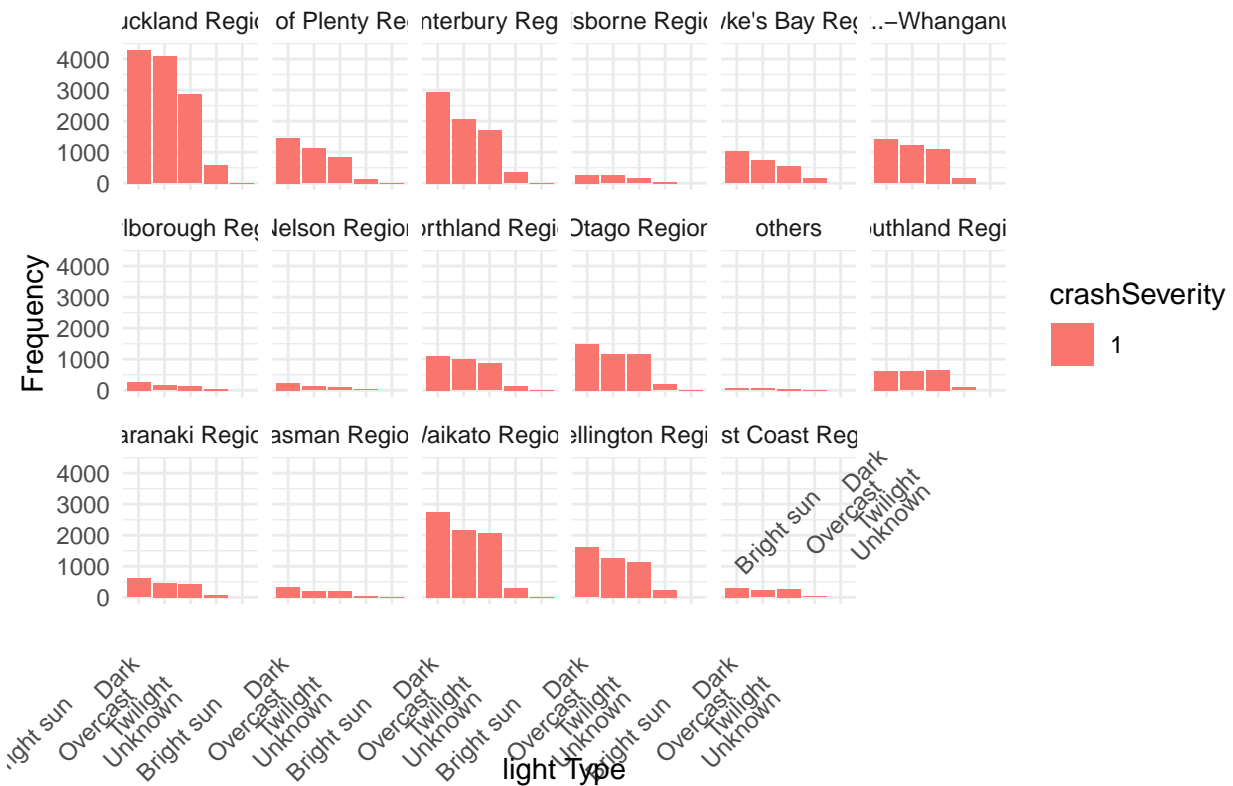


```
## `summarise()` has grouped output by 'region', 'light'. You can override using
## the `.groups` argument.
```

## light Conditions in Region



## light Conditions in Region





## 5. Analytical Plan

Because of the unbalance of the dataset. We decided to equally sample from severe crashes and regular crashes. Because it is catalog dataset, we will use ensemble based on decision tree. Random forest is my prefer to fit the data. We will use 80% of the sample data as training data, and 20% of the sample data as test data. After fit, Finding the most significant predictor in a Random Forest model. That involves analyzing the importance scores of the predictors. Random Forest calculates the importance of each predictor by measuring how much the accuracy of the model decreases when the values of that predictor are randomly permuted while keeping other variables constant. The larger the decrease in accuracy, the more important the predictor is considered.

```
#sample data

# Split data by class
class_data <- split(data, data$crashSeverity)

# Determine desired sample size (e.g., proportionally to the original class distribution)
desired_sample_size <- 5000 # Adjust as needed

# Sample each class
sampled_data <- lapply(class_data, function(class_subset) {
  # Determine the sample size for this class
  class_size <- nrow(class_subset)
  class_sample_size <- min(class_size, desired_sample_size)

  # Sample observations from this class
  sampled_indices <- sample(1:class_size, size = class_sample_size, replace = TRUE)

  # Return the sampled subset
  return(class_subset[sampled_indices, ])
})

# Combine sampled subsets
balanced_data <- do.call(rbind, sampled_data)
```

fit a random forest model

```
#fit a model and test

training_idx <- sample(nrow(balanced_data), nrow(balanced_data)*0.8)
test_idx <- (1:nrow(balanced_data))[-training_idx]

training_data <- balanced_data[training_idx,]
test_data <- balanced_data[test_idx,]

#train_Y <- training_data$crashSeverity


#train_X <- training_data %>% select(-crashSeverity)

test_Y <- test_data$crashSeverity
test_X <- test_data %>% select(-crashSeverity)

library(randomForest)

rf_model <- randomForest(crashSeverity ~ ., data = training_data)
```

```
plot(rf_model)
```



```
predictions <- predict(rf_model, newdata = test_X)
```

```
accuracy <- mean(predictions == test_Y)
```

```
print(accuracy)
```

```
## [1] 0.734
```

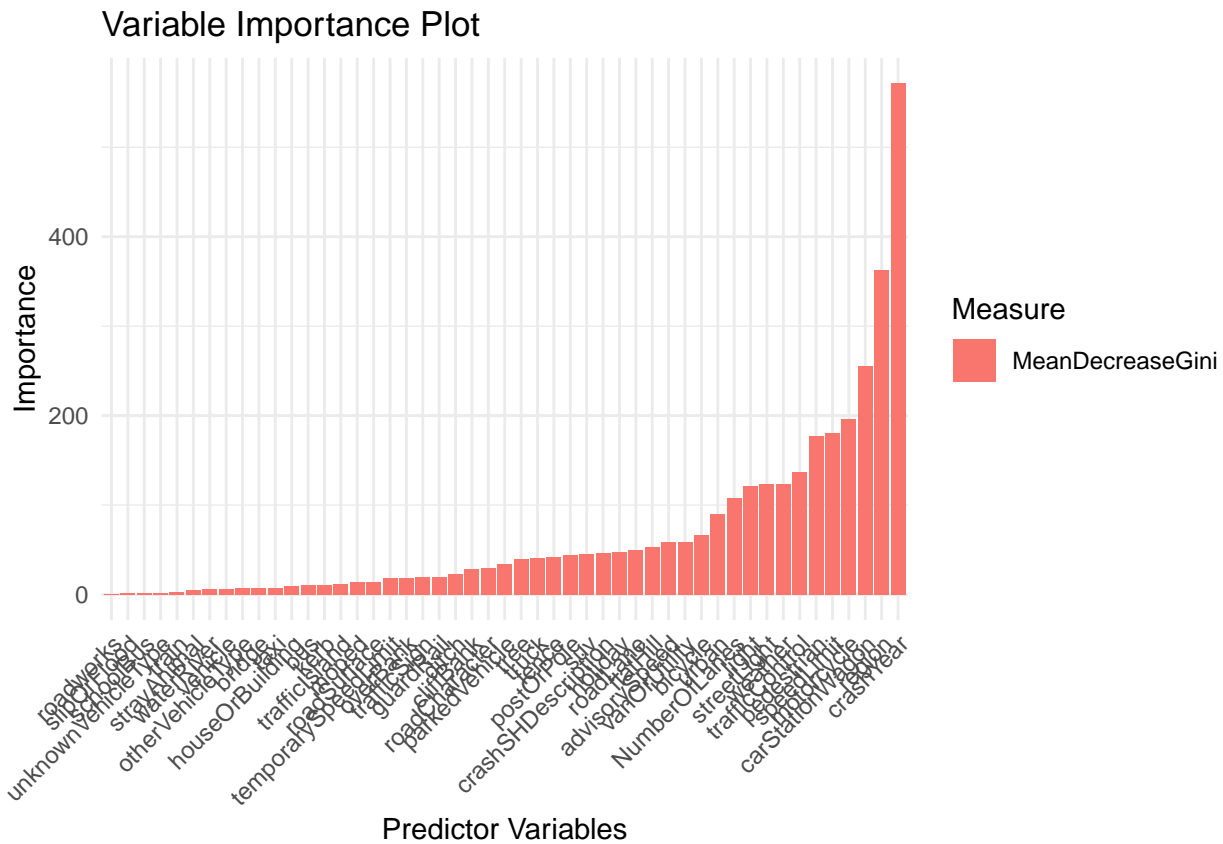
```
knitr::kable(table(predictions, test_Y))
```

	0	1
0	653	181
1	351	815

analysis the importance

	MeanDecreaseGini
crashYear	571.32273
region	362.82414
carStationWagon	254.95824
motorcycle	196.26954
speedLimit	180.78257
pedestrian	177.59802
trafficControl	136.55451
weather	123.78827

	MeanDecreaseGini
streetLight	123.31054
light	120.80554
NumberOfLanes	108.10591
urban	89.57288
bicycle	66.82217
vanOrUtility	58.64939
advisorySpeed	58.21941
flatHill	52.68257
roadLane	50.15274
holiday	47.84270
crashSHDescription	46.19015
suv	45.52350



## 6. Summarise

We fitted a random forest model based on sample data. and the predict accuracy is 0.734, that is quit good. From Sorted importance table and bar char.We can find the most important factors: crashYear 571.3227293 region 362.8241371 carStationWagon 254.9582385 motorcycle 196.2695437 speedLimit 180.7825677 pedestrian 177.5980152 trafficControl 136.5545083 weather 123.7882740 streetLight 123.3105373 light 120.8055409 NumberOfLanes 108.1059101 urban 89.5728839 bicycle 66.8221703 vanOrUtility 58.6493884 advisorySpeed 58.2194140 flatHill 52.6825686 roadLane 50.1527358 holiday 47.8426982 crashSHDescription 46.1901514 suv 45.5235026

##. Appendix:

```

# Path to your plain text file
text_file <- "rcode.txt"

# Read and include the content of the text file
cat(readLines(text_file), sep = "\n")

## library(tidyverse)
## library(xgboost)
##
## #file_path <- file.choose()
## set.seed(230)
## data <- read.csv("../data/Crash_Analysis_System_(CAS)_data.csv", header = TRUE, sep = ",")
##
## dim(data)
##
##
##
## #clean data
##
## columns_to_drop <- c("X","Y","OBJECTID","areaUnitID","crashDirectionDescription","", "crashDistance",
## "crashFinancialYear","fatalCount","debris","meshblockId","northing","easting","objectThrownOrDropped
## "otherObject","phoneBoxEtc","seriousInjuryCount")
##
## data <- select(data, -one_of(columns_to_drop))
##
## na_percentage <- colMeans(is.na(data))
## columns_with_high_na <- names(na_percentage[na_percentage > 0.99])
## #print(columns_with_high_na)
##
## data <- data %>% select(-columns_with_high_na)
##
## #table(data$crashSeverity)
##
## data <- data %>%
##   mutate(crashSeverity = ifelse(crashSeverity == "Fatal Crash" | crashSeverity == "Serious Crash", 1
##                               ifelse(crashSeverity == "Minor Crash" | crashSeverity == "Non-Injury Crash", 0
##                               2))) %>% filter(crashSeverity != 2)
##
## #table(data$crashSeverity)
##
## #table(data$weatherA)
## #table(data$weatherB)
##
## data$weatherA <- ifelse(data$weatherA %in% c("None", "Null"), "Other", data$weatherA)
## data$weatherB<- ifelse(data$weatherB %in% c("None", "Null"), "", data$weatherB)
##
## data <- data %>% unite(weatherA,weatherB,col=weather,sep=" ")
##
##
## table(data$region)
##
## data$region <- ifelse(data$region %in% c("None", "Null"), "Other", data$region)
##
## na_columns <- sapply(data, function(x) any(is.na(x)))

```

```

##
## columns_with_na <- names(data)[na_columns]
## #print(columns_with_na)
##
## data <- data %>%
##   mutate_at(vars(one_of(columns_with_na)), ~replace_na(., 0))
##
##
## processed_data <- data
## processed_data[processed_data == ""] <- "others"
##
## processed_data[] <- lapply(processed_data, function(x) as.factor(x))
## data <- processed_data
##
## rm(processed_data)
##
## #glimpse(data)
##
##
## #explore data
##
## library(ggplot2)
##
##
## #unbalanced target Label
## knitr::kable(table(data$crashSeverity))
##
## ggplot(data, aes(x = crashSeverity)) +
##   geom_bar(fill = "blue") +
##   labs(title = "Histogram of crashSeverity", x = "crashSeverity", y = "Frequency")
##
##
## group_by_region <- data %>% group_by(crashYear, region, crashSeverity) %>% summarise(count=n())
##
## ggplot(group_by_region, aes(x = region, y = count, fill = crashSeverity)) +
##   geom_bar(stat = "identity", position = "dodge") +
##   labs(title = "Comparison of Severe and Regular Values by Region", x = "Region", y = "count") +
##   theme(axis.text.x = element_text(angle = 45, hjust = 2))+
##   facet_wrap(~crashYear, nrow=3)
##
## group_by_region <- group_by_region %>% filter(crashSeverity==1)
##
## ggplot(group_by_region, aes(x = region, y = count, fill = crashSeverity)) +
##   geom_bar(stat = "identity", position = "dodge") +
##   labs(title = "Comparison of Severe and Regular Values by Region", x = "Region", y = "count") +
##   theme(axis.text.x = element_text(angle = 45, hjust = 2))+
##   facet_wrap(~crashYear, nrow=3)
##
## group_by_weather <- data %>% group_by(region, weather, crashSeverity) %>% summarise(count=n())
## ggplot(group_by_weather, aes(x = weather, y = count, fill = crashSeverity)) +
##   geom_bar(stat = "identity", position = "dodge") +
##   labs(title = "Weather Conditions in Region", x = "Weather Type", y = "Frequency") +

```

```

##   theme_minimal() +
##   theme(axis.text.x = element_text(angle = 45, hjust = 2)) + # Rotate x-axis labels if needed
##   facet_wrap(~region,nrow=3)
##
## group_by_weather <- group_by_weather %>% filter(crashSeverity==1)
## ggplot(group_by_weather, aes(x = weather, y = count, fill = crashSeverity)) +
##   geom_bar(stat = "identity", position = "dodge") +
##   labs(title = "Weather Conditions in Region", x = "Weather Type", y = "Frequency") +
##   theme_minimal() +
##   theme(axis.text.x = element_text(angle = 45, hjust = 2)) + # Rotate x-axis labels if needed
##   facet_wrap(~region,nrow=3)
##
## group_by_light <- data %>% group_by(region,light,crashSeverity) %>% summarise(count=n())
## ggplot(group_by_light, aes(x = light, y = count, fill = crashSeverity)) +
##   geom_bar(stat = "identity", position = "dodge") +
##   labs(title = "light Conditions in Region", x = "light Type", y = "Frequency") +
##   theme_minimal() +
##   theme(axis.text.x = element_text(angle = 45, hjust = 2)) + # Rotate x-axis labels if needed
##   facet_wrap(~region,nrow=3)
##
## group_by_light <- group_by_light %>% filter(crashSeverity==1)
## ggplot(group_by_light, aes(x = light, y = count, fill = crashSeverity)) +
##   geom_bar(stat = "identity", position = "dodge") +
##   labs(title = "light Conditions in Region", x = "light Type", y = "Frequency") +
##   theme_minimal() +
##   theme(axis.text.x = element_text(angle = 45, hjust = 2)) + # Rotate x-axis labels if needed
##   facet_wrap(~region,nrow=3)
##
##
## #sample data
##
## # Split data by class
## class_data <- split(data, data$crashSeverity)
##
## # Determine desired sample size (e.g., proportionally to the original class distribution)
## desired_sample_size <- 5000 # Adjust as needed
##
## # Sample each class
## sampled_data <- lapply(class_data, function(class_subset) {
##   # Determine the sample size for this class
##   class_size <- nrow(class_subset)
##   class_sample_size <- min(class_size, desired_sample_size)
##
##   # Sample observations from this class
##   sampled_indices <- sample(1:class_size, size = class_sample_size, replace = TRUE)
##
##   # Return the sampled subset
##   return(class_subset[sampled_indices, ])
## })
##
## # Combine sampled subsets
## balanced_data <- do.call(rbind, sampled_data)
##
##

```

```

##
##
## #fit a model and test
##
## #fit a model and test
##
## training_idx <- sample(nrow(balanced_data), nrow(balanced_data)*0.8)
## test_idx <-(1:nrow(balanced_data))[-training_idx]
##
## training_data <- balanced_data[training_idx,]
## test_data <- balanced_data[test_idx,]
##
## #train_Y <- training_data$crashSeverity
##
## #train_X <- training_data %>% select(-crashSeverity)
##
##
## test_Y <- test_data$crashSeverity
## test_X <- test_data %>% select(-crashSeverity)
##
## library(randomForest)
##
## rf_model <- randomForest(crashSeverity ~ ., data = training_data)
##
## plot(rf_model)
##
## predictions <- predict(rf_model, newdata = test_X)
##
## accuracy <- mean(predictions == test_Y)
##
## print(accuracy)
##
## knitr::kable(table(predictions,test_Y))
##
##
## #analyse the importance
##
## importance_measures <- importance(rf_model)
## sorted_importance <- importance_df[order(importance_df$MeanDecreaseGini, decreasing = TRUE), , drop = FALSE]
##
## knitr::kable(head(sorted_importance,20))
##
##
## importance_df <- as.data.frame(importance_measures)
##
## # Add variable names as a column
## importance_df$Variable <- rownames(importance_df)
##
## # Reshape the data for ggplot
## importance_df_long <- tidyr::gather(importance_df, key = "Measure", value = "Importance", -Variable)
##
## # Create the variable importance plot using ggplot
## ggplot(importance_df_long, aes(x = reorder(Variable, Importance), y = Importance, fill = Measure)) +

```

```
## geom_bar(stat = "identity") +  
## labs(title = "Variable Importance Plot",  
##       x = "Predictor Variables",  
##       y = "Importance",  
##       fill = "Measure") +  
## theme_minimal() +  
## theme(axis.text.x = element_text(angle = 45, hjust = 1))
```